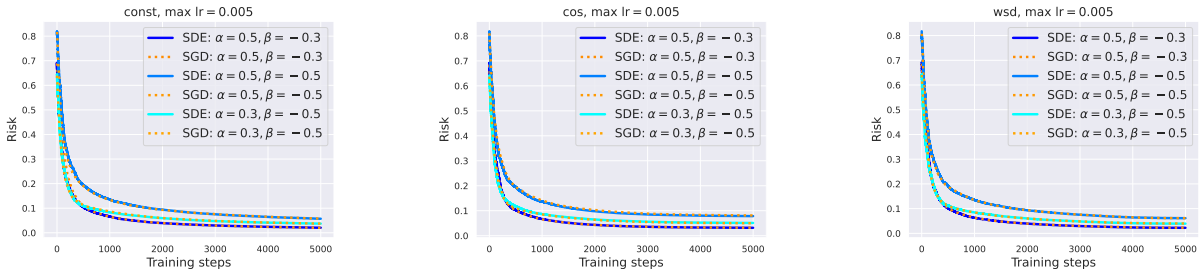

(a) Constant LRS (b) Cosine LRS (c) WSD LRS

Figure 1: SDE and SGD loss comparison under our linear regression setup with different LRS and α, β's

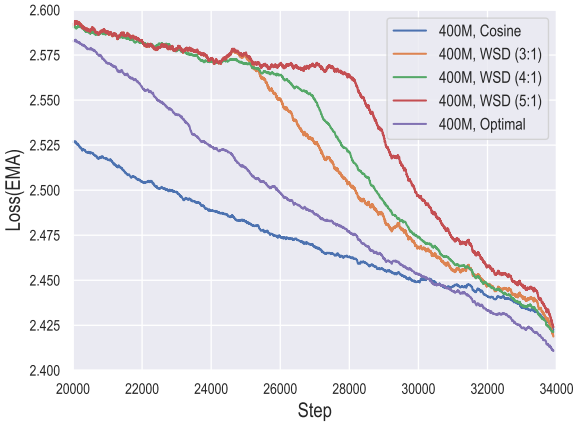

Figure 2: 400M Llama (dense) model, 20B data. WSD with different decaying ratio compared with the optimal

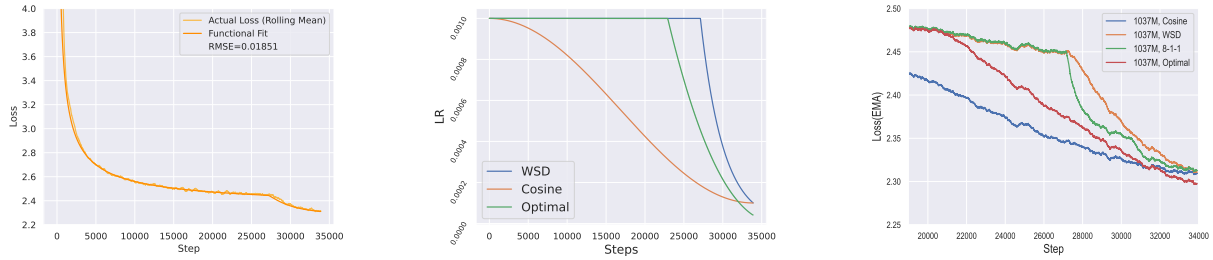

(a) Fitting result (b) Optimal LRS (c) Loss curves

Figure 3: 1B Llama (dense) model, 20B data

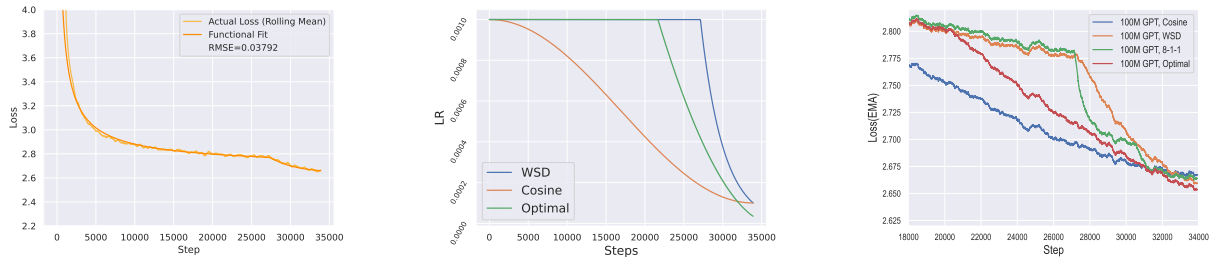

(a) Fitting result (b) Optimal learning rate (c) Loss curves

Figure 4: 100M GPT2 (dense) model, 20B data

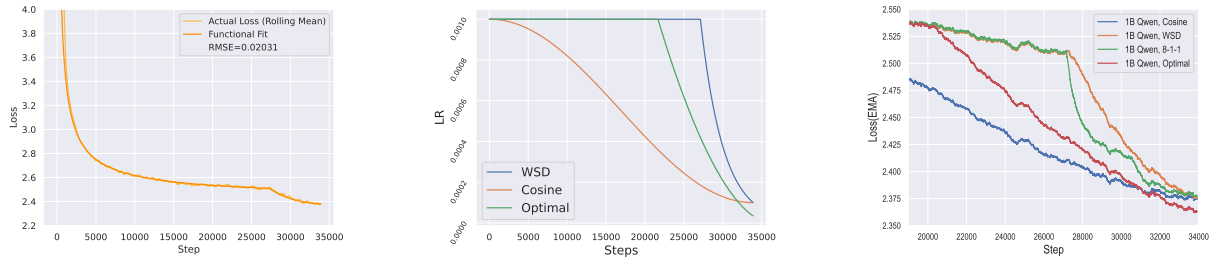

(a) Fitting result (b) Optimal learning rate (c) Loss curves

Figure 5: 1B Qwen (MoE) model, 20B data

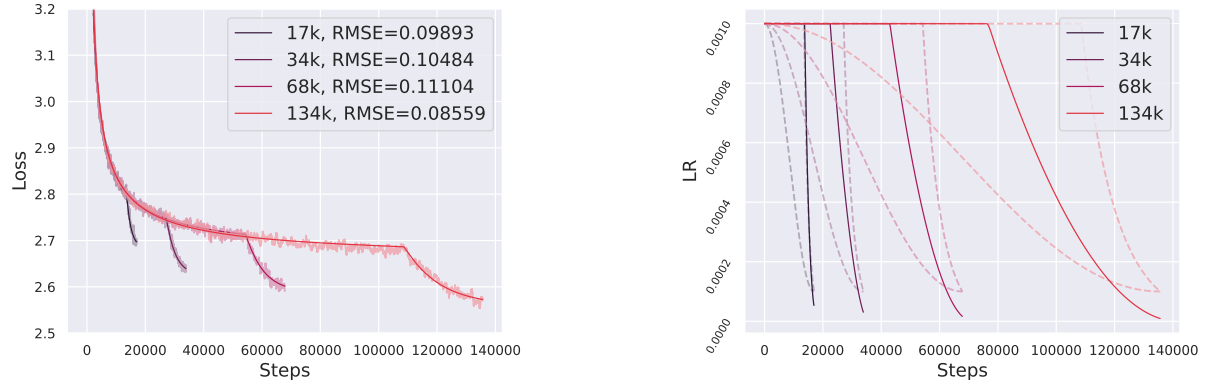

(a) Fitting result (b) Optimal learning rate

Figure 6: Fitted functional scaling laws and optimal LRSs on 100M Llama model with different total training steps. In (b), the solid lines are optimal LRSs, the dashed lines are cosine/WSD LRSs for comparison

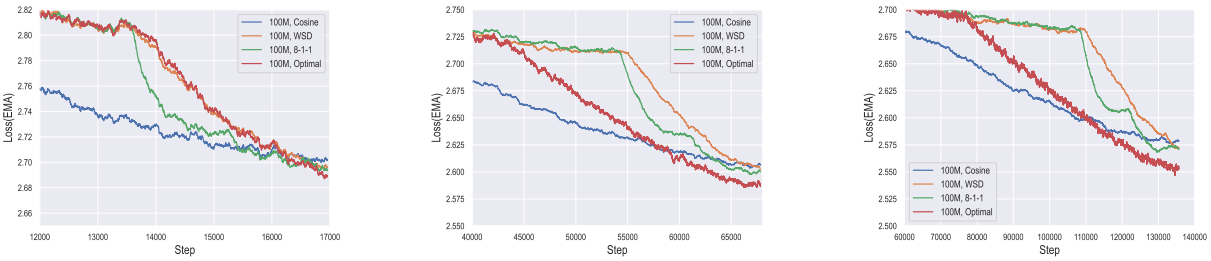

(a) 17k steps (b) 68k steps (c) 134k steps

Figure 7: Loss curves on 100M Llama model with different total training steps

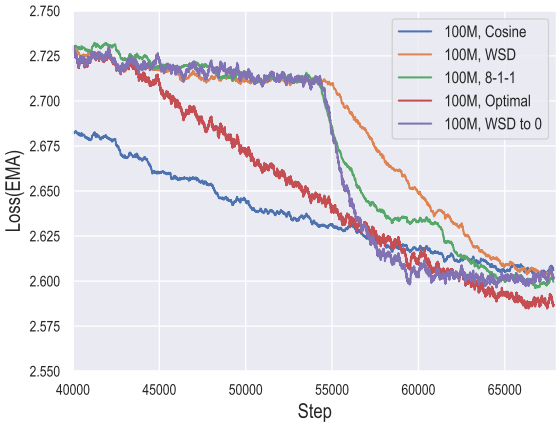

Figure 8: 100M Llama (dense) model, 40B data. WSD decaying to zero does not bring significant loss reduction